Attention Is All You Need

Ashish Vaswani, Noam Shazeer, Niki Parmar

1. Introduction

Recurrent neural networks have been firmly established as state of the art models in sequence modeling. We propose a new simple network architecture, the Transformer, based solely on attention mechanisms.

2. Method

Multi-head attention allows the model to jointly attend to information from different representation subspaces. Scaled dot-product attention weighs queries against keys and returns a weighted sum over values.

3. Conclusion

We achieve state of the art results on translation tasks while cutting training time substantially.